



Distributed Orchestration Approaches for Per-Source FIFO with Retry

This document outlines the problem, requirements, and architecture patterns to ensure strict per-source FIFO processing, retry-until-success, and multi-region/multi-instance scalability.

Problem Statement

Design a system that:

- Processes operations in FIFO order **per sourceId**.
 - **Retries until success** for each message before moving to the next.
 - Supports **multi-region deployment**.
 - Works across **multiple ECS/Fargate instances**.
 - Ensures **only one message per source is processed at a time**.
 - Offers **resiliency, observability, and scalability**.
-

Functional Requirements

- Strict ordering of messages per source.
- Message is processed **only after previous one for the same source is complete**.
- Message should **retry until success**, not go to DLQ.
- Scalable horizontally and regionally.

♥ Non-Functional Requirements

- High availability.
 - Fault tolerance.
 - Low operational overhead.
 - Cloud-native compatibility (AWS services preferred).
-

Approach 1: DynamoDB Only (Polling and Locking Logic)

Architecture Diagram

```
graph TD; A[Client/Event Source] --> B[DynamoDB Global Table]; B --> C[Polling Workers (ECS/Lambda)]
```

↓
Main App with Locking Mechanism

How it Works

- Messages are written to a DynamoDB table partitioned by `sourceId`.
- Polling workers continuously scan for unprocessed messages in FIFO order.
- Use conditional writes (e.g., `UpdateItem` with `ConditionExpression`) to acquire a lock on the oldest unprocessed message.
- Retry logic is implemented in code.
- Message is only marked processed if successful.

Advantages

- No need for additional AWS services (e.g., SQS).
- Complete control over retry, ordering, and locking.
- Multi-region support via global tables.

Disadvantages

- Requires precise design of locking, polling, and failure recovery.
- Higher complexity in custom orchestration code.
- Potential read amplification and cost if not optimized.

Best For

- Teams comfortable writing custom orchestration.
- Scenarios with strong FIFO and retry consistency requirements.

Approach 2: DynamoDB + SQS FIFO (Direct Worker Model)

Architecture Diagram

```
graph TD;
    A[Client/Event Source] --> B[DynamoDB Global Table];
    B --> C["SQS FIFO Queue (MessageGroupId = sourceId)"];
    C --> D[ECS/Fargate Workers];
    D --> E[Main App];
```

How it Works

- Operations are written to DynamoDB.
- Message pushed to SQS FIFO queue with `MessageGroupId = sourceId`.
- Multiple workers consume messages.
- FIFO maintained per source.
- Retry handled by SQS (visibility timeout).

Advantages

- Strict per-source FIFO via `MessageGroupId`.
- Auto-scaling across ECS instances.
- Retry handled by infrastructure.
- Easy to operate and scale.

Disadvantages

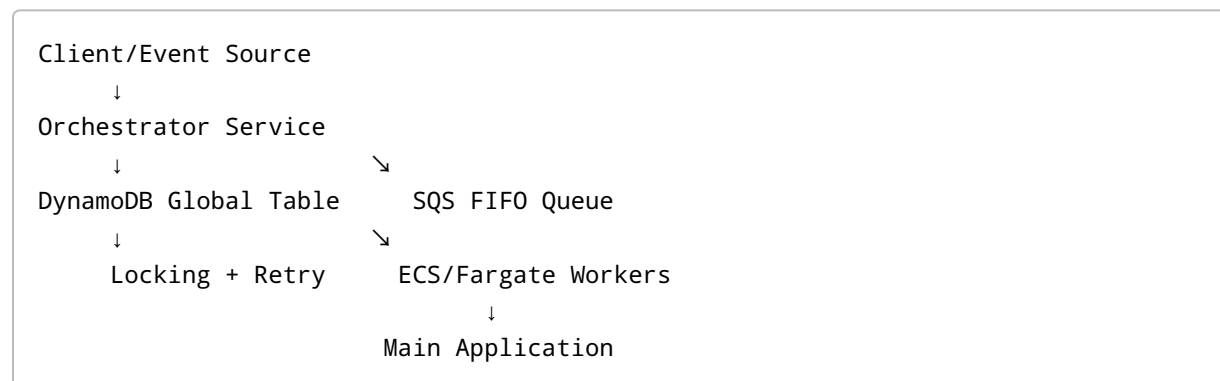
- Message retry behavior is timer-based (not conditional).
- No global routing unless you design it.
- Need to manage stuck messages with watchdog/DLQ.

Best For

- High concurrency, multi-source event pipelines.
- Simpler orchestration, fully infra-driven retry.

Approach 3: Orchestrator Service + Main Application (Separation of Concerns)

Architecture Diagram



How it Works

- Orchestrator accepts writes, pushes them to DynamoDB & SQS.
- Orchestrator polls and controls FIFO + retry logic.

- Main application is stateless, only processes calls.
- Orchestrator enforces: "next message processed only if previous is successful."

Advantages

- Full control of ordering and retry behavior.
- Clean decoupling between business logic and orchestration.
- Easy to extend (e.g., throttling, alerting, metrics).
- Supports regional queues + routing.

Disadvantages

- Requires additional orchestration code and deployments.
- Slightly more operational complexity.

Best For

- Enterprise-grade pipelines.
- Multi-region, multi-tenant architectures.
- Fine-grained control over ordering and retries.

Other Notable Approaches (Short Summary)

Step Functions + DynamoDB Streams

- Use DynamoDB Streams to trigger Step Functions state machines per source.
- State machine enforces FIFO by maintaining execution state.
- Complex for large-scale concurrency and limited flexibility with custom retry.

Kafka with Partitions by SourceId

- Use Kafka with topic partitions keyed by `sourceId` to ensure per-source FIFO.
- Retry can be implemented with a separate retry topic.
- Needs self-managed infra or MSK.

SNS + Lambda FIFO Emulation

- SNS to fan out events; use `sourceId` hash to route to designated Lambda or SQS FIFO.
- Requires careful mapping and routing logic.
- Works better with moderate scale.

Aurora + Advisory Locks

- Use a relational DB (like Aurora) with advisory locks for source-level locking.
- Suitable where SQL querying and ordering is a priority.
- Higher latency and complexity at scale.

Feature Comparison Table

| Feature | DynamoDB Only | SQS FIFO Direct | Orchestrator Service |
|-------------------------------|---------------|-----------------|----------------------|
| Strict per-source FIFO | Custom | Native via SQS | Full control |
| Retry until success | Custom | Via SQS | Custom retry |
| Multi-region routing | Global Table | with design | Flexible |
| Horizontal scalability | With effort | Yes | Yes |
| Orchestration observability | Minimal | Minimal | Full logs/metrics |
| Dev/Operational simplicity | Complex | Moderate | Complex |
| Fine-grained ordering control | Custom logic | By group | Full |

Recommendations

- **DynamoDB Only:** When full control is needed and you're comfortable building orchestration from scratch.
- **SQS FIFO:** For high-scale FIFO pipelines with infra-managed simplicity.
- **Orchestrator Service:** For enterprise-grade, region-aware, retry-until-success flows with rich observability and control.

Let me know if you'd like this exported to PDF, with diagrams or editable format.